

Before we start

About me

- Petra Kaferle Devisschere
- Data Scientist @ Adaltas
 - Working with clients on PoCs and production-destined projects
 - Teaching: Spark, Python, Git, and for the first time MLOps
 - Technical articles: <https://www.adaltas.com/en/articles/>
- DSTI Data science S19
 - Background in biochemistry with a lot of data analysis and visualization


About the course

- NEW => please, give feedback
- 4 days, 2 modules per day
- Material for the labs: <https://github.com/adaltas/dsti-mlops-2021-spring>
- Focus on understanding the concept, not on coding

Do interact and ask questions

Motivation

- A lot of MLOps skills implement good practice (versioning, testing...) => increase the **quality** and **reproducibility** of work
- **High in-demand skills!**
- To avoid the following situations:



IT WORKS
ON MY MACHINE

1. Introduction to MLOps

Bringing ML projects to production

Why is it important?

- Only models in production bring added value to the business
- Data Scientists are not final users

But...

- 80% of the models are never deployed
- Developing a model takes weeks/months, deploying takes months/year(s)

Data science landscape

Data science is much more than just machine learning!

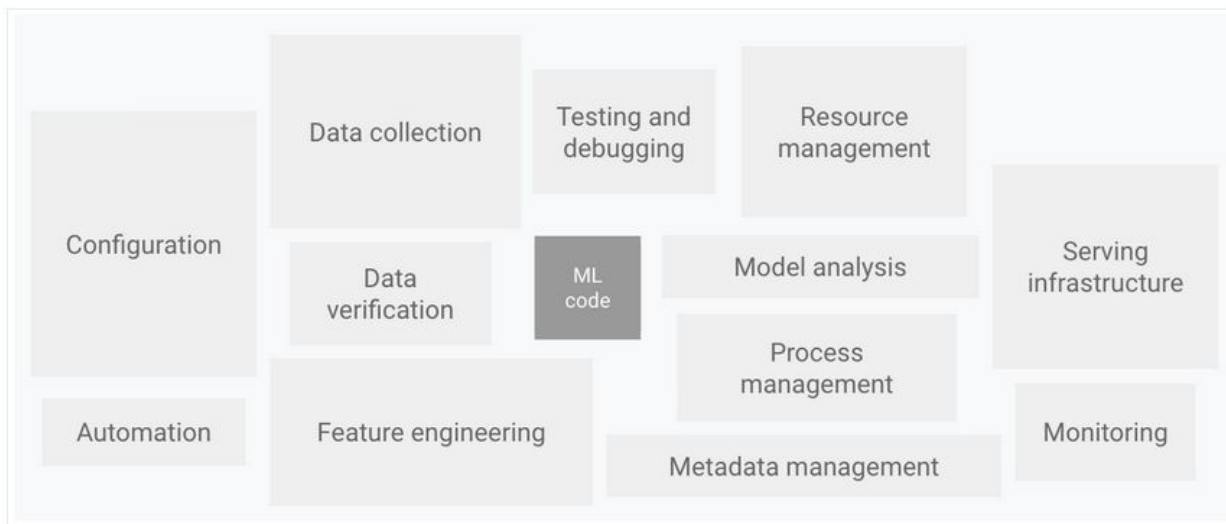


Figure 1. Elements for ML systems. Adapted from [Hidden Technical Debt in Machine Learning Systems](#).

ML in research vs. production

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference, low latency
Data	Static	Constantly shifting
Fairness	Good to have (sadly)	Important
Interpretability*	Good to have	Important

Challenges of deploying ML


Table 1: All considerations, issues and concerns explored in this study. Each is assigned to the stage and step of the deployment workflow where it is commonly encountered.

Deployment Stage	Deployment Step	Considerations, Issues and Concerns
Data management	Data collection	Data discovery
	Data preprocessing	Data dispersion Data cleaning
	Data augmentation	Labeling of large volumes of data Access to experts Lack of high-variance data
	Data analysis	Data profiling
Model learning	Model selection	Model complexity Resource-constrained environments Interpretability of the model
	Training	Computational cost Environmental impact
	Hyper-parameter selection	Resource-heavy techniques Hardware-aware optimization
Model verification	Requirement encoding	Performance metrics Business driven metrics
	Formal verification	Regulatory frameworks
	Test-based verification	Simulation-based testing
Model deployment	Integration	Operational support Reuse of code and models Software engineering anti-patterns Mixed team dynamics
	Monitoring	Feedback loops Outlier detection Custom design tooling
	Updating	Concept drift Continuous delivery
Cross-cutting aspects	Ethics	Country-level regulations Focus on technical solution only Aggravation of biases Authorship Decision making
	End users' trust	Involvement of end users User experience Explainability score
	Security	Data poisoning Model stealing Model inversion

Software is being deployed all the time, so what's the big deal?

Let's start with analogy in software development => **DevOps**

But what is DevOps?



- **It's a culture**
- Automation
- Quick delivery of results

DevOps Principles (Software development)

Version Control. Developers submit code changes to a central repository several times a day. Prior to submitting code to the master repository (master branch), all code must be verified. To facilitate collaboration, other developers can track changes.

Continuous Integration. Members of the development team integrate their code in a shared repository, several times a day. Each developer segments the work into small, manageable chunks of code and detects potential merge conflicts and bugs quicker.

Continuous Delivery. As the code is continuously integrated, it is also consistently delivered to the end-user. Smaller contributions allow faster update releases, which is a crucial factor for customer satisfaction.

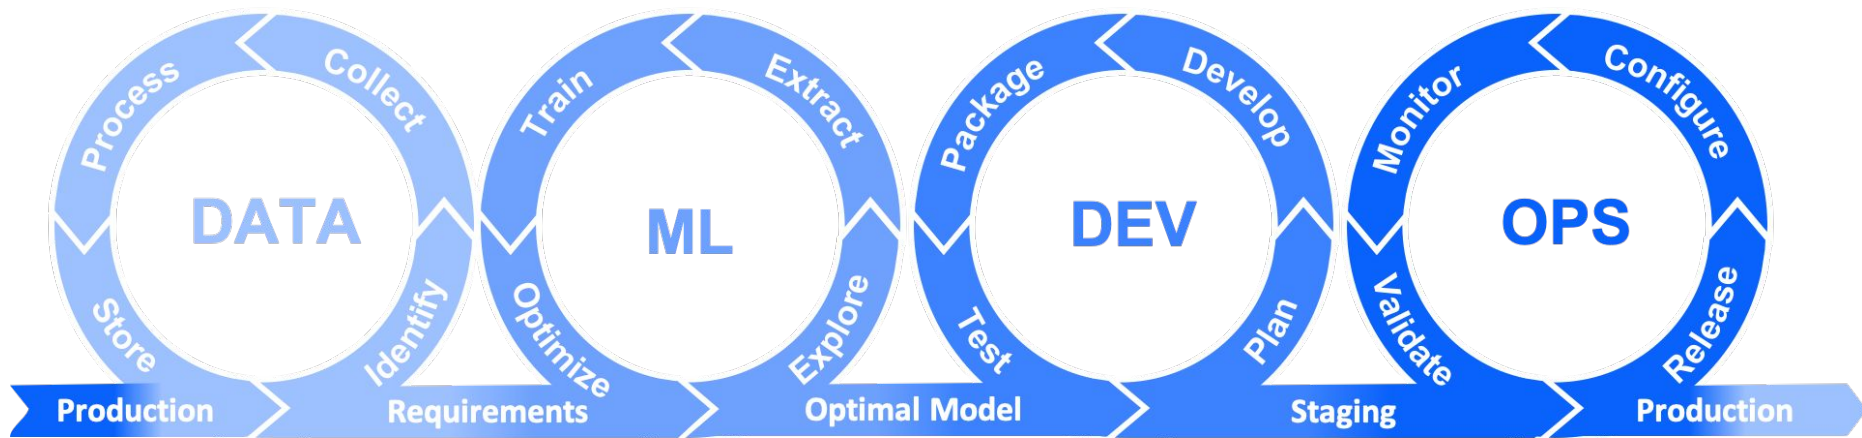
Continuous Deployment. A big part of DevOps is automating processes to speed up production. Continuous deployment involves automating releases of minor updates that do not pose a substantial threat to the existing architecture.

Continuous Testing. Such a strategy involves testing as much as possible in every step of development. Automated tests give valuable feedback and a risk assessment of the process at hand.

Continuous Operations. The DevOps team is always working on upgrading software with small but frequent releases. That is why DevOps requires constant monitoring of performance. Its main goal is to prevent downtime and availability issues during code release.

Collaboration. One of the main goals of DevOps is to foster collaboration and feedback sharing. Development and Operations need to proactively communicate and share feedback to maintain an efficient DevOps pipeline.

MLOps = DevOps for ML



MLOps Principles :

replace `code` with `code + model + data` => MLOps

Version Control. Developers submit code changes to a central repository several times a day. Prior to submitting code to the master repository (master branch), all code must be verified. To facilitate collaboration, other developers can track changes.

Continuous Integration. Members of the development team integrate their code in a shared repository, several times a day. Each developer segments the work into small, manageable chunks of code and detects potential merge conflicts and bugs quicker.

Continuous Delivery. As the code is continuously integrated, it is also consistently delivered to the end-user. Smaller contributions allow faster update releases, which is a crucial factor for customer satisfaction.

Continuous Deployment. A big part of DevOps is automating processes to speed up production. Continuous deployment involves automating releases of minor updates that do not pose a substantial threat to the existing architecture.

Continuous Testing. Such a strategy involves testing as much as possible in every step of development. Automated tests give valuable feedback and a risk assessment of the process at hand.

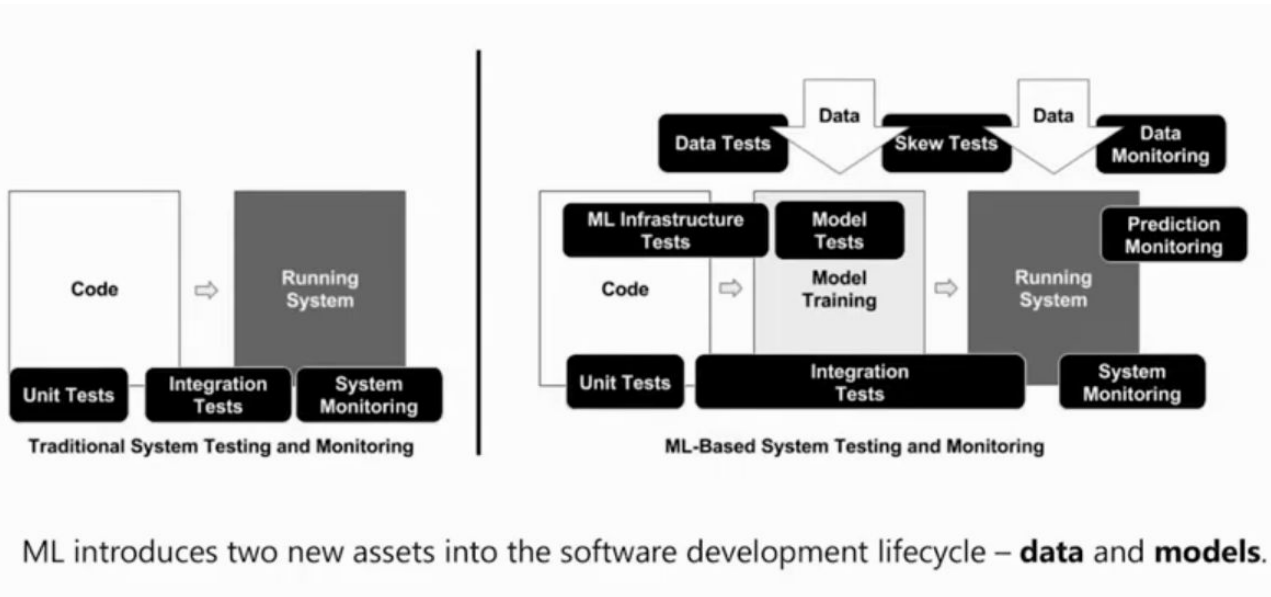
➡ **Continuous Operations.** The DevOps team is always working on upgrading software with small but frequent releases. That is why DevOps requires constant monitoring of performance. Its main goal is to prevent downtime and availability issues during code release.

Collaboration. One of the main goals of DevOps is to foster collaboration and feedback sharing. Development and Operations need to proactively communicate and share feedback to maintain an efficient DevOps pipeline.

➡ **Model feedback.** Saving incoming data and predictions once a model is deployed

DevOps vs MLOps

We deploy:
- code (app)

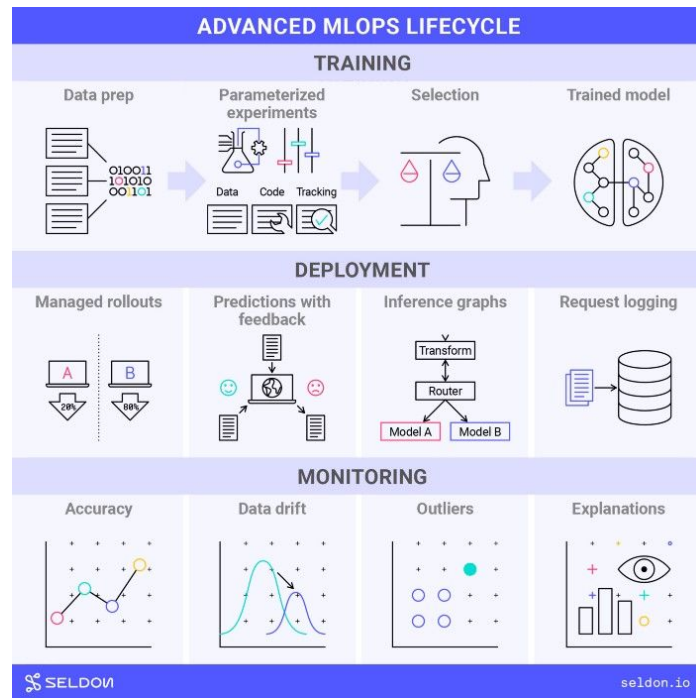
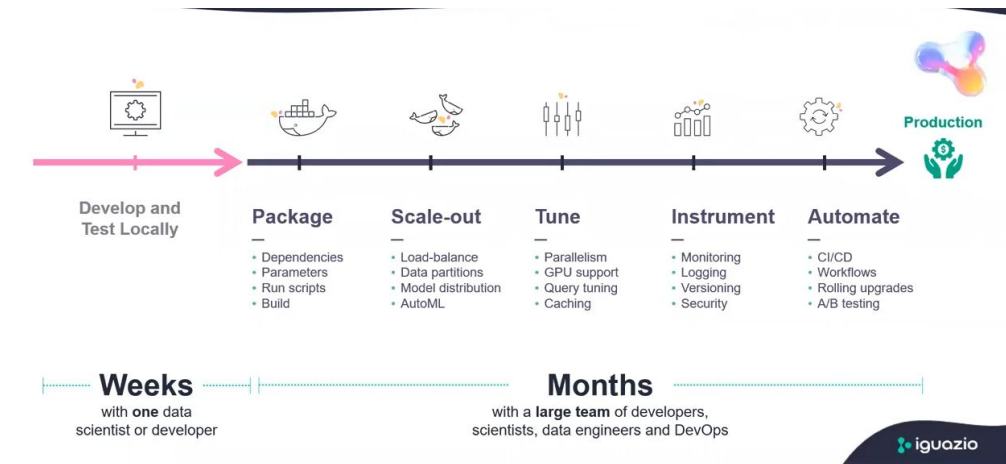


We deploy:

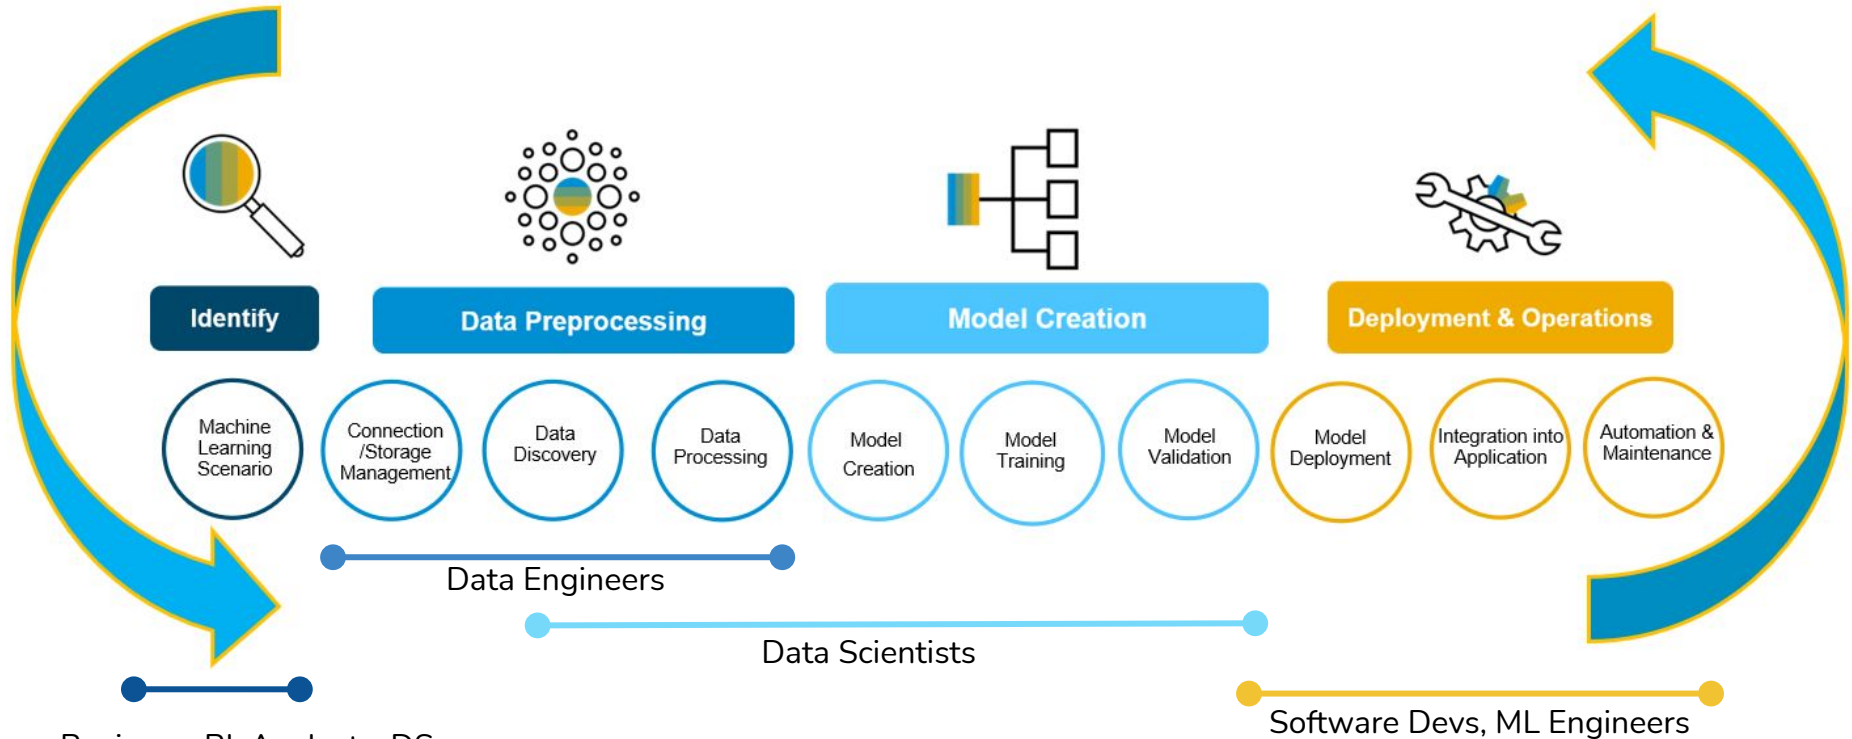
- code:
 - + data pre-processing
 - + feature engineering
- model
- sometimes data
 - + aggregates

=> **pipeline**

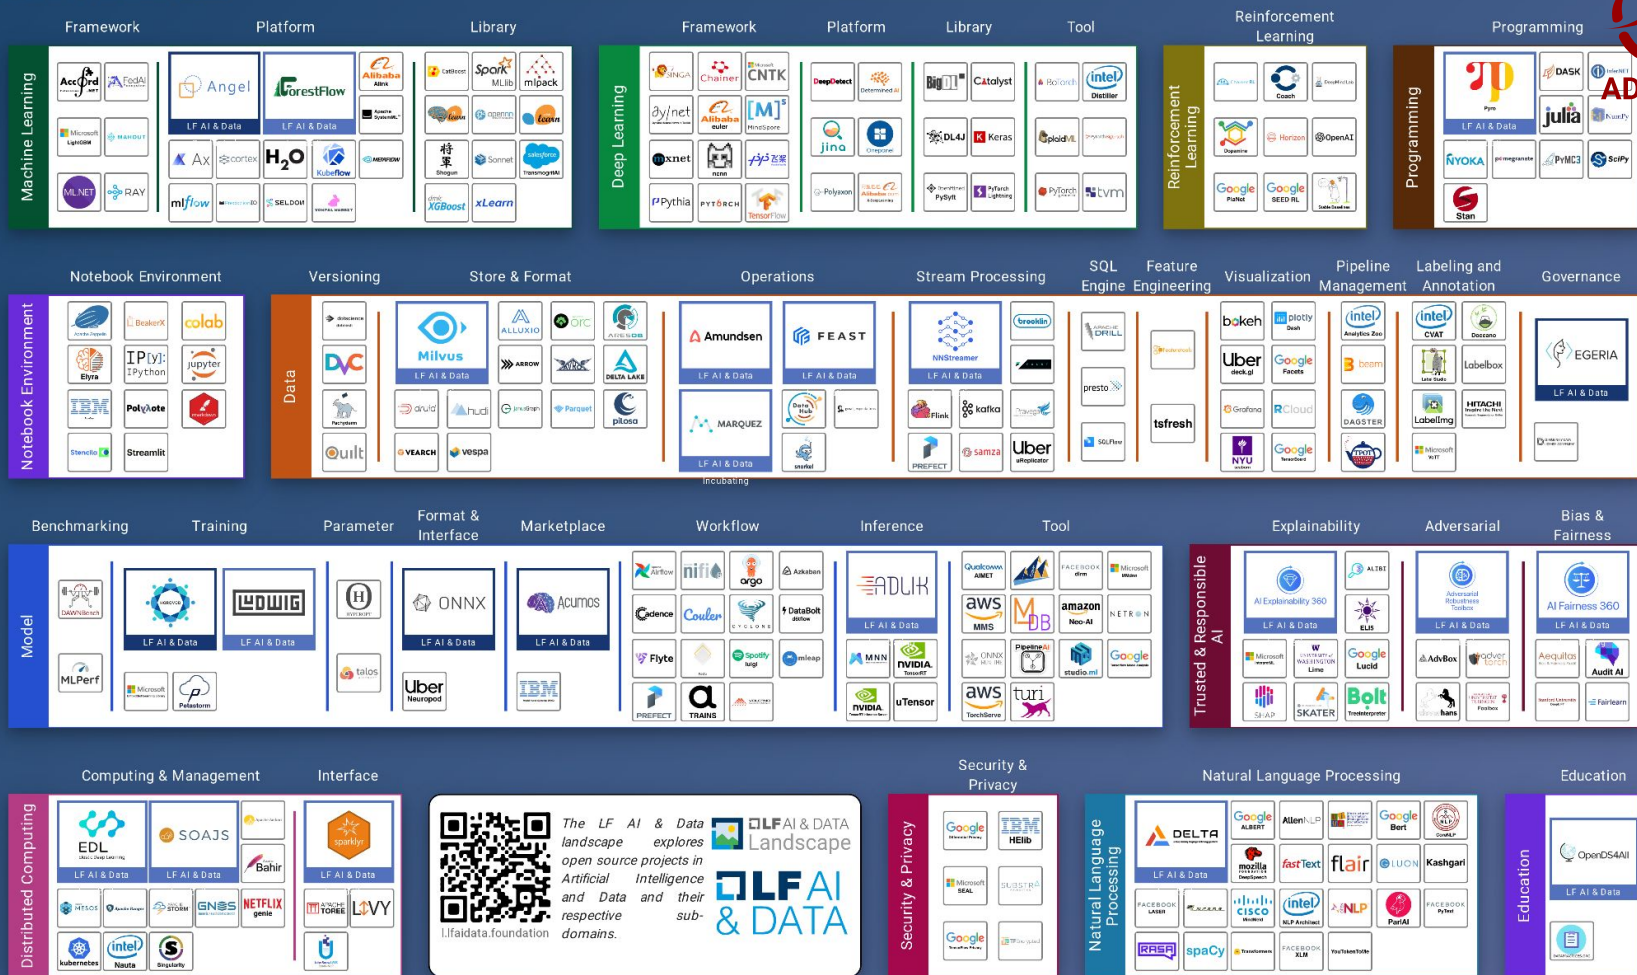
Different tasks (and definitions) of MLOps



ML Lifecycle and Roles



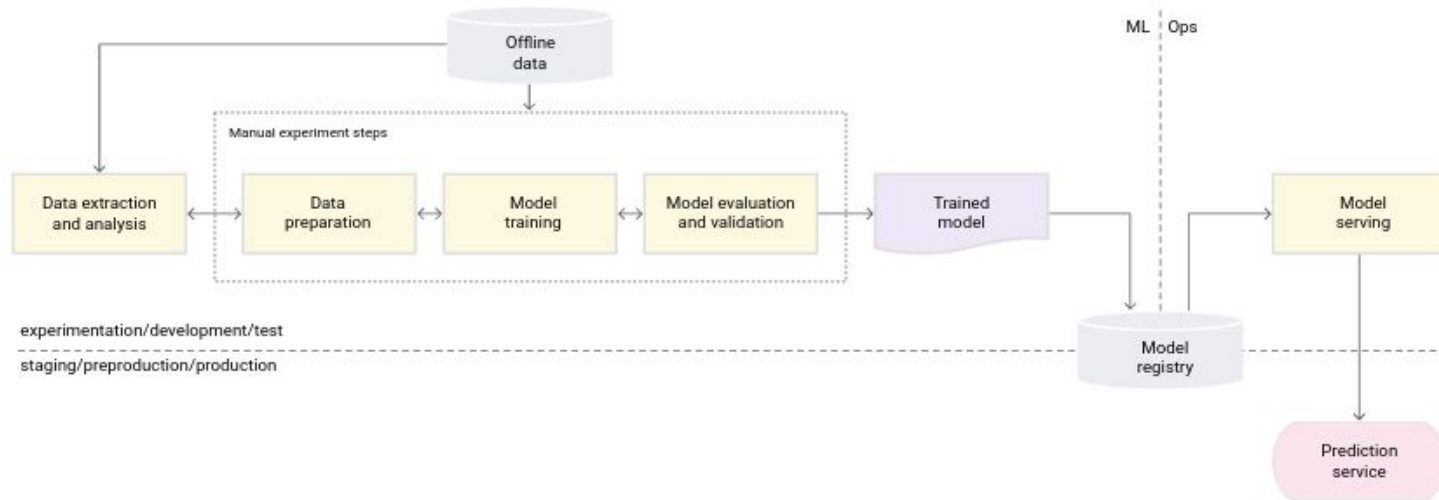
THINK ABOUT THE TOOLS !!!



Maturity levels of MLOps systems (according to Google)

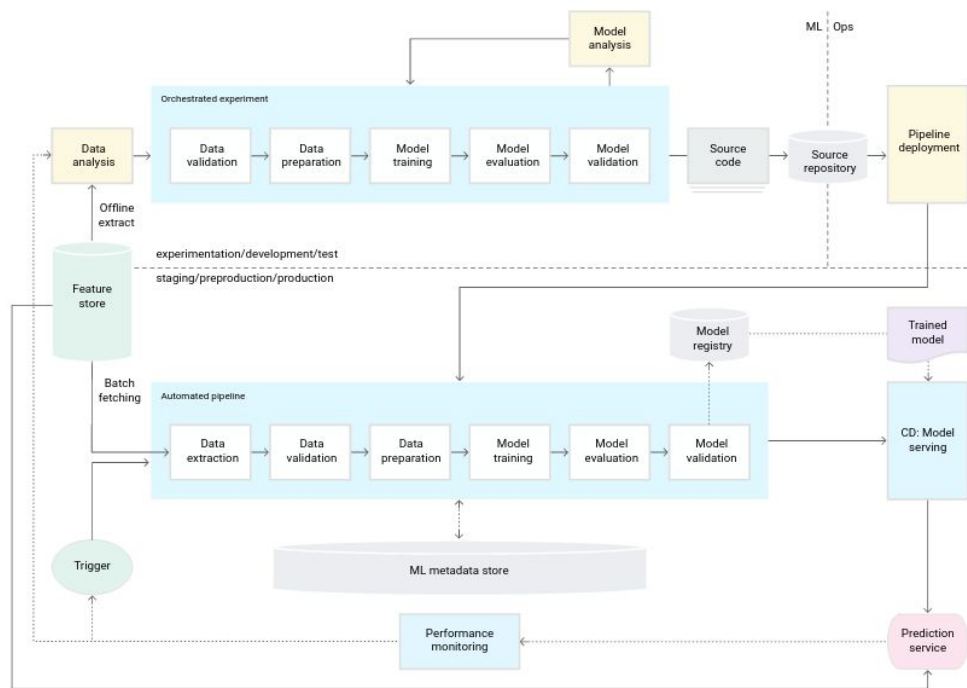
- Level 0 - manual
- Level 1- ML pipeline automation
- Level 2 - CI/CD pipeline automation

MLOps level 0: manual process



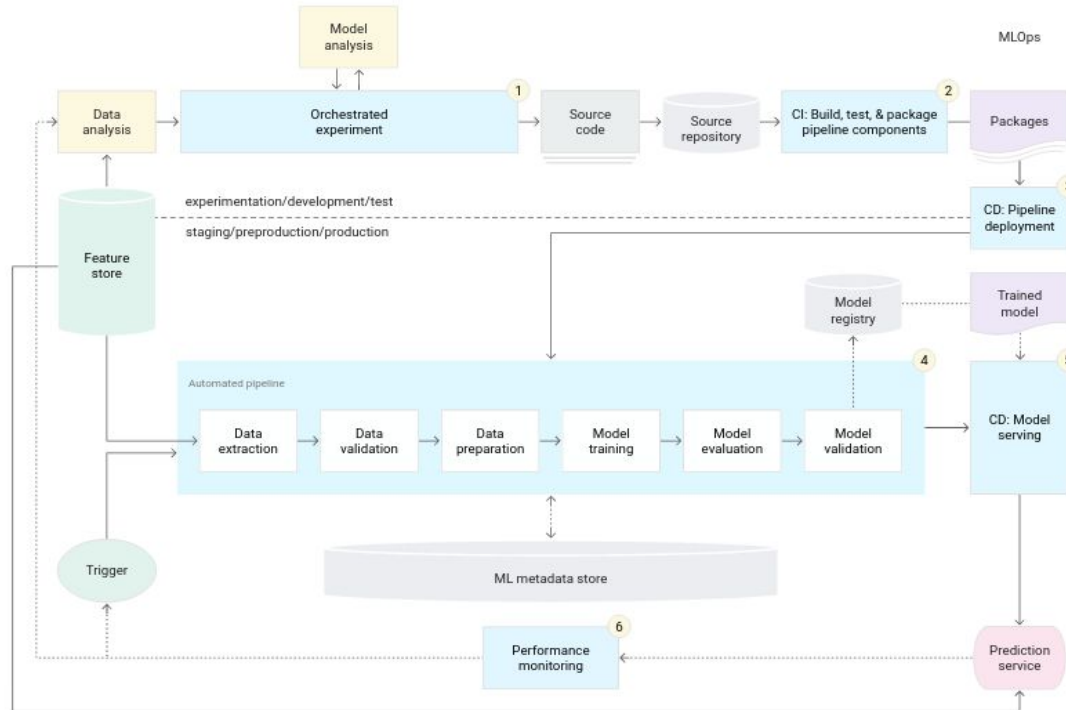
- Data scientists and Ops teams are separated
- Only the model is served (not the whole pipeline)
- Model is not frequently updated
- No monitoring

MLOps level 1: ML pipeline automation



- The goal is CT in production when **new data** is available
- We need to deploy the whole pipeline (data+model)
- The steps need to be automated and orchestrated
- Code is reproducible and modular
- Need of additional components:
 - Triggers - to trigger the re-training
 - Feature store - stores pre-computed features that model uses during serving
 - Metadata store - execution logs

MLOps level 2: CI/CD pipeline automation



- Automatically build, test, and deploy the new pipeline components to the target environment
- Deployment:
 - Automated - to the test env.
 - Semi-automated - to pre-production (triggered by merge)
 - Manual - to production

How does it work in real life?

- Depends a lot on the skills and experience of the group
- Depends a lot on the number of models to deploy/maintain
- Depends on how often we need to retrain the model

- Generally:
 - Starting with the simplest system
 - Gradual improvements: automation of critical steps, adding tests, adding model monitoring...